



Advanced Time Series Analysis and Forecasting

Seminar 6: State space models and Bayesian filtering



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AI: Artificial Intelligence
- ▣ APF: Auxiliary Particle Filter
- ▣ AR: AutoRegressive (model); in event studies: Abnormal Return
- ▣ ARIMA: AutoRegressive Integrated Moving Average (model)
- ▣ ARMA: AutoRegressive Moving Average (model)
- ▣ BN: Beveridge–Nelson (decomposition)
- ▣ BNR: Banca Națională a României (the National Bank of Romania)
- ▣ CEE: Central and Eastern Europe
- ▣ DFM: Dynamic Factor Model
- ▣ DK: Durbin–Koopman (simulation smoother, 2002)
- ▣ ESS: Effective Sample Size
- ▣ FFBS: Forward Filtering, Backward Sampling
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GDP: Gross Domestic Product
- ▣ HICP: Harmonised Index of Consumer Prices
- ▣ IG: inverse gamma (distribution)
- ▣ IMA: Integrated Moving Average (model)
- ▣ KSC: Kim–Shephard–Chib (1998), mixture sampler for stochastic volatility
- ▣ LR: Likelihood Ratio (test)
- ▣ MA: Moving Average
- ▣ ML: Machine Learning
- ▣ MNZ: Morley–Nelson–Zivot (2003)
- ▣ PMMH: Particle Marginal Metropolis–Hastings
- ▣ QML: Quasi-Maximum Likelihood
- ▣ SCA: Seasonally and Calendar Adjusted
- ▣ SV: Stochastic Volatility

Acronyms used in this material (2/2)

- **TSA: Time Series Analysis**
- **TVP: Time-Varying Parameter(s)**
- **UC: Unobserved Components (model)**
- **UC-SV: Unobserved Components model with Stochastic Volatility (Stock–Watson 2007)**
- **US: United States**
- **VAT: Value Added Tax**

Today's question and route

- **Question:** how do we estimate what we cannot observe (a level, a trend, a volatility, a gap), and how do we tell a real signal from the noise of the estimator?
 - ▶ this seminar comes **before** Lecture 6: the section “What you need today” gives every definition the tasks use
- **Route**
 - ▶ Part A: the exact diffuse filter by hand, the concentrated likelihood, backward sampling, the Gibbs block, SV moments, one particle-filter step, the APF, Beveridge–Nelson
 - ▶ Part B: Romanian trend growth; TVP regressions for Hungary and Poland; SV of EUR/RON by Gibbs and by the particle filter; the real-time Romanian output gap
 - ▶ Part C: is Romanian trend inflation anchored? An AI answer to audit
- Notebook for today: [open the seminar notebook in Google Colab](#); each task names its notebook section

Exercise map

Task	Question	Type	Model
A1, A2	the exact diffuse Kalman filter by hand; the concentrated likelihood	Solved, Proposed	A1
A3, A4	one backward-sampling step; the inverse-gamma Gibbs block	Solved, Proposed	A3
A5, A6	moments of the SV model; the QML measurement equation	Solved, Proposed	A5
A7, A8	one bootstrap particle-filter step; auxiliary weights with an outlier	Solved, Proposed	A7
A9, A10	the Beveridge–Nelson cycle; local level, IMA(1,1) and BN	Solved, Proposed	A9
B1, B2	Romanian trend growth; TVP regressions for Hungary and Poland	Solved, Proposed	B1
B3, B4	SV of EUR/RON: Gibbs, then the particle filter	Solved, Proposed	B3
B5	the Romanian output gap in real time	Proposed	B1, A10
C1, C2	is Romanian trend inflation anchored?; what is wrong in an AI answer?	Proposed	B3, A1–A10

- ▣ **[Solved]**: full solution in the slides and in the notebook, a model to follow; **[Proposed]**: you solve it, following the model

Data used

Series	Source	Frequency	Period
Romania: real GDP (SCA)	Eurostat namq_10_gdp	quarterly	2000–2026
HICP: Romania, Hungary, Poland, euro area	Eurostat prc_hicp_minr	monthly	2001–2026
EUR/RON reference rate	BNR	daily	2016–2026

- In the notebook: `kalman_filter`, `kalman_smoother`, `sim_smoother_dk`, `ffbs`, `sv_gibbs`, `pf_sv`, `fit_llt`, `fit_tvp`, `fit_smooth_trend`, `ucsv_gibbs`; no account or key is needed

What you need today (1/4): the filter and the diffuse start

- Model: $y_t = Z_t \alpha_t + \varepsilon_t$, $\varepsilon_t \sim N(0, H)$; $\alpha_{t+1} = T \alpha_t + R \eta_t$, $\eta_t \sim N(0, Q)$ (TSA, Chapter 10)
- Filter: $v_t = y_t - Z_t a_t$, $F_t = Z_t P_t Z_t' + H$, $K_t = T P_t Z_t' / F_t$, $a_{t+1} = T a_t + K_t v_t$, $P_{t+1} = T P_t (T - K_t Z_t')' + R Q R'$; filtered: $a_{t|t} = a_t + P_t Z_t' v_t / F_t$
- Diffuse start: $P_1 = \kappa P_\infty + P_*$, $\kappa \rightarrow \infty$; while $F_{\infty,t} = Z_t P_{\infty,t} Z_t' > 0$: $K^{(0)} = T P_\infty Z' / F_\infty$, $K^{(1)} = T(P_* Z' - P_\infty Z' F_* / F_\infty) / F_\infty$
 - ▶ $a_{t+1} = T a_t + K^{(0)} v_t$, $P_{\infty,t+1} = T P_\infty L^{(0)'}$, $P_{*,t+1} = T P_\infty L^{(1)'} + T P_* L^{(0)'} + R Q R'$, $L^{(0)} = T - K^{(0)} Z$, $L^{(1)} = -K^{(1)} Z$
- Diffuse log-likelihood: the diffuse steps contribute $-\frac{1}{2}(\ln 2\pi + \ln F_{\infty,t})$; the others $-\frac{1}{2}(\ln 2\pi + \ln F_t + v_t^2 / F_t)$
- Scale concentrated out: if H, Q are proportional to σ^2 , then $\hat{\sigma}^2 = \frac{1}{n-d} \sum_{t>d} v_t^2 / F_t^*$

What you need today (2/4): sampling the states

- FFBS (Carter–Kohn): draw $\alpha_n \sim N(a_n|n, P_n|n)$; then $\alpha_t|\alpha_{t+1} \sim N(a_t|t + G_t(\alpha_{t+1} - Ta_{t|t}), P_t|t - G_tTP_{t|t}),$
 $G_t = P_{t|t}T'(TP_{t|t}T' + RQR')^{-1}$
- Durbin–Koopman: simulate (α^+, y^+) from the model, smooth $y - y^+$ once: $\tilde{\alpha} = \alpha^+ + \hat{\alpha}(y - y^+)$
- Gibbs block for a variance with prior $IG(a_0, b_0)$ and m shocks e_t : $\sigma^2|e \sim IG(a_0 + m/2, b_0 + \sum e_t^2/2)$, mean $b/(a - 1)$
- Inefficiency factor $1 + 2 \sum_k \rho_k$: how many correlated draws are worth one independent draw
- Pile-up: the ML estimate of a variance can be exactly zero even when the true value is positive; LR tests at the boundary need a bootstrap

What you need today (3/4): stochastic volatility and particles

- ▣ SV: $y_t = e^{h_t/2} \epsilon_t$, $h_t = \mu + \phi(h_{t-1} - \mu) + \sigma_\eta \eta_t$; $\sigma_h^2 = \sigma_\eta^2 / (1 - \phi^2)$; kurtosis $3e^{\sigma_h^2}$; $\text{Corr}(y_t^2, y_{t-k}^2) = (e^{\sigma_h^2 \phi^k} - 1) / (3e^{\sigma_h^2} - 1)$
- ▣ $\ln y_t^2 = h_t + \ln \epsilon_t^2$; $\mathbb{E} \ln \epsilon^2 = -1.2704$, $\text{Var} \ln \epsilon^2 = \pi^2/2$; KSC: a seven-normal mixture for $\ln \epsilon^2$, so h is Gaussian given the indicators
- ▣ Bootstrap particle filter: propagate $x_t^{(i)} \sim f(\cdot | x_{t-1}^{(i)})$; weight $w^{(i)} = g(y_t | x_t^{(i)})$; $\hat{p}(y_t | Y_{t-1}) = \frac{1}{N} \sum w^{(i)}$; ESS = $1 / \sum (W^{(i)})^2$; resample
- ▣ Systematic resampling: $u_i = (U + i - 1)/N$; particle j is copied for every u_i in $(C_{j-1}, C_j]$, C_j = cumulative normalised weights
- ▣ Auxiliary filter (Pitt–Shephard): first-stage weights $\propto g(y_t | \mu_t^{(i)})$, $\mu_t^{(i)} = \mathbb{E}(x_t | x_{t-1}^{(i)})$; then propagate and correct

What you need today (4/4): trend and cycle

- Beveridge–Nelson: $\tau_t = y_t + \sum_{j \geq 1} \mathbb{E}_t(\Delta y_{t+j} - \mu)$; for $\Delta y_t - \mu = \phi(\Delta y_{t-1} - \mu) + e_t$: $c_t = y_t - \tau_t = -\frac{\phi}{1-\phi}(\Delta y_t - \mu)$
- Local level $y_t = \mu_t + \varepsilon_t$, $\mu_{t+1} = \mu_t + \eta_t$, $q = \sigma_\eta^2 / \sigma_\varepsilon^2$: Δy_t is MA(1), $\Delta y_t = u_t + \theta u_{t-1}$ with $\theta = -(1 + q/2) + \sqrt{q + q^2/4}$
- Local linear trend: $\mu_{t+1} = \mu_t + \nu_t + \xi_t$, $\nu_{t+1} = \nu_t + \zeta_t$; $4\nu_t$ is the annualised trend growth of a quarterly log series
- Smooth-trend UC: $y_t = \tau_t + c_t$, $\Delta^2 \tau_{t+1} = \zeta_t$, c_t AR(2); filtered gap $c_{t|t}$ (real time) against smoothed gap $\hat{c}_t$ (full sample)
- UC-SV (Stock–Watson): $\pi_t = \tau_t + \sigma_{\eta,t} \eta_t$, $\tau_t = \tau_{t-1} + \sigma_{\varepsilon,t} \varepsilon_t$, log-variances as random walks with variance γ

A1: the exact diffuse filter by hand [Solved]

Task

- ▣ Local level, $\sigma_\varepsilon^2 = 1$, $\sigma_\eta^2 = 0.5$, μ_1 diffuse ($a_1 = 0$, $P_{\infty,1} = 1$, $P_{*,1} = 0$); data $y = (2.0; 3.0; 2.5)$.
- ▣ 1. Run the diffuse step at $t = 1$: give F_∞ , $K^{(0)}$, $K^{(1)}$, a_2 , $P_{\infty,2}$, $P_{*,2}$.
- ▣ 2. Run the ordinary filter for $t = 2, 3$: v_t , F_t , K_t , a_{t+1} , P_{t+1} .
- ▣ 3. Compute the diffuse log-likelihood and say why y_1 carries no information about the variances.
- ▣ Report: one diffuse step, two ordinary steps, one number, one sentence.

Solution

- ▣ 1. $F_\infty = 1$, $F_* = 0 + 1 = 1$; $K^{(0)} = 1$, $K^{(1)} = 0 - 1 \cdot 1/1 = -1$;
 $a_2 = 0 + 1 \cdot 2 = 2.000$; $P_{\infty,2} = 1(1 - 1) = 0$;
 $P_{*,2} = 1 \cdot 1 + 0 + 0.5 = 1.500$
- ▣ 2. $t = 2$: $v = 1.000$, $F = 2.500$, $K = 0.600$, $a_3 = 2.600$,
 $P_3 = 1.5(1 - 0.6) + 0.5 = 1.100$; $t = 3$: $v = -0.100$,
 $F = 2.100$, $K = 0.524$, $a_4 = 2.548$, $P_4 = 1.024$
- ▣ 3. Terms -0.919 , -1.577 , -1.292 ; $\ln L_d = -3.788$. The first term is $-\frac{1}{2} \ln 2\pi$ whatever the variances: y_1 only fixes the level

A2: the concentrated likelihood [Proposed]

Solution discussed in class

Task

- ▣ The data of A1, now with $\sigma_\eta^2 = 0.5\sigma_\varepsilon^2$ and σ_ε^2 unknown. Model: A1.
- ▣ 1. Run the filter with $\sigma_\varepsilon^2 = 1$ and write $F_t = \sigma_\varepsilon^2 F_t^*$.
- ▣ 2. Derive $\hat{\sigma}_\varepsilon^2$ by maximising the diffuse log-likelihood and compute it.
- ▣ 3. Compute the concentrated log-likelihood (without the diffuse term) and state how many parameters remain for the numerical search.
- ▣ Report: two ratios, one estimate, one number, one sentence.

A3: one backward-sampling step [Solved]

Task

- ▣ Continue A1: after y_2 the filtered moments are $a_{2|2}$, $P_{2|2}$; suppose the draw for the next period is $\alpha_3 = 3.0$.
- ▣ 1. Compute $a_{2|2}$ and $P_{2|2}$ from A1.
- ▣ 2. Compute G_2 , the conditional mean m_2 and variance S_2 of $\alpha_2 | \alpha_3, y_1, y_2$.
- ▣ 3. Explain why y_3 is not needed in this conditional distribution.
- ▣ Report: two moments, three numbers, one sentence.

Solution

- ▣ 1. $a_{2|2} = 2 + 0.6 \cdot 1 = 2.600$, $P_{2|2} = 1.5 - 1.5^2/2.5 = 0.600$
- ▣ 2. $G_2 = P_{2|2}/(P_{2|2} + \sigma_\eta^2) = 0.545$;
 $m_2 = 2.600 + 0.545(3.0 - 2.600) = 2.818$;
 $S_2 = P_{2|2}(1 - G_2) = 0.273$
- ▣ 3. Markov property: given α_3 , the future observations add nothing about α_2 ; their information already entered through the draw of α_3

A4: the inverse-gamma Gibbs block [Proposed]

Solution discussed in class

Task

- ▣ A local level Gibbs sampler, $n = 50$; current state draw gives $\sum_t (y_t - \mu_t)^2 = 40$ and $\sum_t (\Delta \mu_t)^2 = 6.2$; priors $IG(2.5, 0.25)$ for both variances. Model: A3.
- ▣ 1. Write the two conditional posteriors with their parameters.
- ▣ 2. Compute their means and the implied q .
- ▣ 3. Explain why this sampler can never return $\sigma_\eta^2 = 0$, unlike ML.
- ▣ Report: two distributions, three numbers, one sentence.

A5: moments of the SV model [Solved]

Task

- ▣ SV model with $\phi = 0.98$, $\sigma_\eta = 0.15$, Gaussian ϵ_t .
- ▣ 1. Compute σ_h^2 and the kurtosis of y_t .
- ▣ 2. Compute $\text{Corr}(y_t^2, y_{t-k}^2)$ for $k = 1, 10, 50$ and the half-life of a volatility shock.
- ▣ 3. Explain the difference from GARCH(1,1) with the same persistence.
- ▣ Report: two numbers, four numbers, one sentence.

Solution

- ▣ 1. $\sigma_h^2 = 0.0225/0.0396 = 0.568$; kurtosis $3e^{0.568} = 5.30$
- ▣ 2. 0.173, 0.138, 0.054: slow decay from a low start; half-life $\ln 0.5 / \ln 0.98 = 34.3$ days
- ▣ 3. GARCH variance is known one step ahead; in SV it has its own shock, so excess kurtosis arises even with ϕ close to 1 and the likelihood needs an integral

A6: the QML measurement equation [Proposed]

Solution discussed in class

Task

- ▣ The SV model of A5, written as $\ln y_t^2 = h_t + \xi_t$. Model: A5.
- ▣ 1. Give the mean and variance of ξ_t and the signal-to-noise ratio $\sigma_h^2 / \text{Var} \xi_t$.
- ▣ 2. Compute $\text{Corr}(\ln y_t^2, \ln y_{t-k}^2)$ for $k = 1, 10$ and name the ARMA model of $\ln y_t^2$.
- ▣ 3. Explain why the Kalman filter on $\ln y_t^2$ is consistent but inefficient.
- ▣ Report: three numbers, two numbers and a model, one sentence.

A7: one step of the bootstrap particle filter [Solved]

Task

- ▣ SV model, $N = 4$ particles after propagation:
 $h^{(i)} = (-1; 0; 0.5; 1.5)$; observation $y_t = 2.0$; systematic resampling with $U = 0.3$.
- ▣ 1. Compute $w^{(i)} = N(y_t; 0, e^{h^{(i)}})$ and the normalised weights.
- ▣ 2. Compute the ESS and the likelihood increment $\hat{p}(y_t | Y_{t-1})$.
- ▣ 3. Resample: which particles survive, and how many times?
- ▣ Report: eight numbers, two numbers, one list.

Solution

- ▣ 1. $w = (0.0029; 0.0540; 0.0924; 0.1206)$;
 $W = (0.011; 0.200; 0.342; 0.447)$
- ▣ 2. ESS = 2.80 of 4; $\hat{p} = \frac{1}{4} \sum w = 0.0675$, $\ln \hat{p} = -2.696$
- ▣ 3. $u = (0.075; 0.325; 0.575; 0.825)$ against
 $C = (0.011; 0.211; 0.553; 1)$: particles 2, 3, 4, 4; the
 low-volatility particle dies, the high-volatility one is copied twice

A8: auxiliary weights with an outlier [Proposed]

Solution discussed in class

Task

- ▣ Local level, $\sigma_\varepsilon^2 = 0.5$, $\sigma_\eta^2 = 0.25$; equally weighted particles $x_{t-1}^{(i)} = (0; 1; 2; 3)$; outlier $y_t = 4.0$. Model: A7.
- ▣ 1. Compute the normalised first-stage weights $\propto N(y_t; x_{t-1}^{(i)}, \sigma_\varepsilon^2)$ and their ESS.
- ▣ 2. Compute the fully adapted weights $\propto N(y_t; x_{t-1}^{(i)}, \sigma_\varepsilon^2 + \sigma_\eta^2)$ and compare.
- ▣ 3. Explain what the bootstrap filter would do with the same outlier.
- ▣ Report: four weights and an ESS, four weights, one sentence.

A9: the Beveridge–Nelson cycle of an ARIMA(1,1,0) [Solved]

Task

- ▣ Quarterly growth: $\Delta y_t - 0.8 = 0.35(\Delta y_{t-1} - 0.8) + e_t$; this quarter $\Delta y_t = 2.0$.
- ▣ 1. Derive $\sum_{j \geq 1} \mathbb{E}_t(\Delta y_{t+j} - \mu)$ for an AR(1).
- ▣ 2. Compute the BN cycle c_t and the long-run multiplier of e_t on the trend.
- ▣ 3. Explain the sign of the cycle after a strong quarter.
- ▣ Report: one formula, two numbers, one sentence.

Solution

- ▣ 1. $\mathbb{E}_t(\Delta y_{t+j} - \mu) = \phi^j(\Delta y_t - \mu)$, sum = $\frac{\phi}{1-\phi}(\Delta y_t - \mu)$
- ▣ 2. $\frac{\phi}{1-\phi} = 0.538$; $c_t = -0.538 \cdot 1.2 = -0.646$; multiplier $1/(1-\phi) = 1.538$
- ▣ 3. With positive growth persistence, a strong quarter raises expected future growth: the trend jumps by more than output, so output is *below* trend

A10: local level, IMA(1,1) and Beveridge–Nelson [Proposed]

Solution discussed in class


Task

- ▣ Local level with $\sigma_\varepsilon^2 = 1$, $q = 0.5$. Model: A9.
- ▣ 1. Find the MA coefficient θ of Δy_t .
- ▣ 2. Solve the steady-state Riccati equation $P = P - P^2/(P + 1) + q$ and give the gain $k = P/(P + 1)$.
- ▣ 3. Show that the BN trend of the IMA(1,1), $y_t + \theta u_t$, equals the steady-state filtered level $a_t|_t$.
- ▣ Report: one number, two numbers, one derivation.

B1: Romanian trend growth [Solved]

- ▣ **Question:** has Romania's trend growth fallen since the 2000s, once the uncertainty of the estimate is taken into account?
- ▣ **Inputs:** 100×log real GDP (Eurostat, SCA), 2000Q1–2026Q2, $T = 106$; 2020Q2–Q4 treated as missing
- ▣ **Tasks**
 1. Estimate the local linear trend by exact diffuse ML with the numpy filter and with statsmodels; compare the parameters and the log-likelihoods.
 2. Smooth the slope and report the annualised trend growth $4\hat{\nu}_t$ with a 90% band in 2005, 2012, 2019 and in the last quarter.
 3. Check whether any variance is at the boundary.
 4. Interpretation: is the fall in trend growth statistically meaningful?
- ▣ **Report:** two parameter sets, four growth rates with bands, one sentence
- ▣ **Notebook:** section B1

B1: solution [Solved]



- ▣ **numpy:** $\sigma_{\varepsilon}^2 = 0.6013$, $\sigma_{\xi}^2 = 3.9577$, $\sigma_{\zeta}^2 = 0.0039$, $\ln L_d = -231.08$; **statsmodels:** 0.6011, 3.9581, 0.0039, -229.24 (the difference in $\ln L$ is $\ln 2\pi$, two diffuse steps)
- ▣ **Trend growth:** 4.1% (s.d. 1.1) in 2005, 3.1% (1.0) in 2012, 2.8% (1.1) in 2019, 2.2% (1.4) at the end
- ▣ **Interpretation:** the slope variance is small, so the trend growth moves slowly; the bands overlap widely, so a fall is suggested but not established; level shocks absorb the 2009 recession


B2: TVP regressions for Hungary and Poland [Proposed]

- **Question:** has the co-movement of Hungarian and Polish inflation with euro-area inflation changed over time?
- **Inputs:** quarterly HICP inflation (% a.r.), 2005Q4 onward; euro-area inflation net of its seasonal pattern; the model of B1 extended to $y_t = c_t + b_t x_t + D_t' \gamma + \varepsilon_t$; model: B1
- **Tasks**
 1. Estimate the TVP regression by exact diffuse ML (five diffuse states) for each country.
 2. Report $\hat{\sigma}_b^2$ and the smoothed b_t with standard errors in 2008Q3, 2019Q4 and 2022Q4.
 3. Compute the LR test of $\sigma_b^2 = 0$ and its $\frac{1}{2} \chi_1^2$ p -value; say why this p -value is only indicative.
 4. Interpretation: is the slope higher than one, and does it differ from Romania (Lecture 6)?
- **Report:** a table of two rows, six slopes with standard errors, two tests, one sentence
- **Notebook:** section B2

B3: stochastic volatility of EUR/RON [Solved]

- **Question:** how persistent is the volatility of the leu against the euro, and how does SV compare with GARCH?
- **Inputs:** daily log returns of the BNR EUR/RON reference rate, %, demeaned, 2016–2026, $T = 2686$; KSC Gibbs sampler with $c = 10^{-3}\text{Var}(y)$
- **Tasks**
 1. Run the KSC Gibbs sampler and report posterior means and 90% intervals of μ, ϕ, σ_η with inefficiency factors.
 2. Estimate a Gaussian GARCH(1,1) on the same data.
 3. Plot the posterior mean of $\exp(h_t/2)$ against the GARCH volatility.
 4. Interpretation: what does the half-life of volatility say about a managed exchange rate?
- **Report:** a table of three parameters, three GARCH numbers, one chart, one sentence
- **Notebook:** section B3

B3: solution [Solved]



- ▣ $\phi = 0.957$ [0.944; 0.969], $\sigma_\eta = 0.541$ [0.488; 0.597], $\mu = -5.628$ [-6.032; -5.210]; inefficiency factors 18, 42, 1 (20 000 draws)
- ▣ GARCH(1,1): $\alpha = 0.225$, $\beta = 0.753$; daily s.d. 0.121%, kurtosis 27.8
- ▣ Interpretation: half-life of a log-volatility shock about 16 days; quiet periods with rare bursts are what a managed float produces; the high kurtosis says that Gaussian ϵ_t is not enough (B4)

B4: the particle filter for the EUR/RON SV model [Proposed]

- **Question:** is SV better than GARCH for EUR/RON on the likelihood scale, and is the particle likelihood reliable here?
- **Inputs:** the data and the posterior means of B3; bootstrap particle filter with systematic resampling; model: B3 and A7
- **Tasks**
 1. Run the filter with $N = 500, 2000, 10\,000$ particles, 30 times each, and report the s.d. of $\ln \hat{L}$.
 2. Report the minimum and median ESS and the date of the minimum.
 3. Compare $\ln \hat{L}$ (large N) with the GARCH maximum.
 4. Interpretation: what does the ESS collapse tell you about the model?
- **Report:** three standard deviations, two ESS values and a date, two log-likelihoods, one sentence
- **Notebook:** section B4

B5: the Romanian output gap in real time [Proposed]

- **Question:** how much does the end-of-sample estimate of the Romanian output gap change when later data arrive?
- **Inputs:** the data of B1; smooth-trend UC with AR(2) cycle (`fit_smooth_trend`); quasi-real time: re-estimate on data up to each quarter from 2012Q4 (parameters every 4 quarters) and keep the filtered gap; model: B1 and A10
- **Tasks**
 1. Compute the series of real-time (filtered, end-of-sample) gaps and the final smoothed gap.
 2. Compute the Orphanides–van Norden statistics: correlation, noise-to-signal ratio (s.d. of revisions over s.d. of the final gap), share of opposite signs.
 3. Report the largest revision.
 4. Interpretation: can monetary policy rely on the current estimate of the gap?
- **Report:** four statistics, one number, one sentence
- **Notebook:** section B5

C1: is Romanian trend inflation anchored? [Proposed]

- ▣ **Question:** is the trend of Romanian inflation inside the BNR target band ($2.5\% \pm 1$ pp), and how robust is the answer?
- ▣ **Inputs:** quarterly HICP inflation (% a.r.); UC-SV of Stock and Watson (2007) with quarterly seasonal dummies (ucsv_gibbs); model: B3
- ▣ **Tasks**
 1. Write a five-line pre-registration: hypothesis, data, the anchoring statistic, the robustness grid, what would falsify the claim.
 2. Estimate the model for start dates 2001 and 2006 and for $\gamma = 0.1$ and 0.2 ; report the trend in the last quarter with a 68% band.
 3. Compute the posterior probability that the trend lies in 1.5–3.5% in the last quarter and in 2019Q4.
 4. Interpretation: what can and cannot be concluded, and how would you turn this into a project?
- ▣ **Report:** a five-line plan, a table of four rows, a project plan
- ▣ **Notebook:** section C1

C2: audit an AI answer [Proposed]

- ▣ A student asked an AI assistant to summarise state space estimation. The answer:
- ▣ (a) The Kalman filter is useless if the errors are not Gaussian.
- ▣ (b) For a random-walk state, set $P_1 = 10^{10}$ and keep all n terms of the likelihood.
- ▣ (c) An ML estimate of the level variance equal to zero proves that the level is constant.
- ▣ (d) SV and GARCH are the same model, since both make volatility persistent.
- ▣ (e) The particle-filter likelihood is unbiased, so its logarithm is unbiased too.
- ▣ (f) The Durbin-Koopman simulation smoother needs only one state-smoother pass per draw.
- ▣ (g) With an AR(1) cycle, the correlation between trend and cycle shocks can be estimated.
- ▣ Tasks
 - ▶ 1. For each statement, say whether it is correct; if not, give the correct statement and, where possible, a number from today's tasks (notebook section C2).
 - ▶ 2. Report: a list of seven verdicts with one line of justification each.

Key takeaways

- ▣ The exact diffuse filter lets the first observations fix the nonstationary states; they carry no information about the variances
- ▣ Sampling the states is the Kalman machinery run backwards (FFBS) or once more on simulated data (DK)
- ▣ SV adds a shock to the variance; particle filters estimate its likelihood, unbiased in levels only
- ▣ Trend growth, gaps and anchored trends are estimates with wide bands: report them, and report real-time values
- ▣ An AI answer is a draft: check the initialisation, the boundary, the log-likelihood bias and the identification

After the seminar

- Lecture 6 develops each topic of today: the exact diffuse filter, the likelihood at the boundary, the three samplers, KSC, particle filters and PMMH, TVP, DFM, MNZ and UC-SV
- Try the [Proposed] tasks in the notebook
- C1 can grow into a team project: trend inflation and anchoring in Central and Eastern Europe
- Reading: [Durbin and Koopman \(2012\)](#), Ch. 4–5; [Kim, Shephard and Chib \(1998\)](#); [Morley, Nelson and Zivot \(2003\)](#); [Stock and Watson \(2007\)](#)
- **The seminar is for practice and is not graded; the solutions of [Proposed] tasks are discussed in class**

References (1/2)

- Beveridge, S., & Nelson, C. R. (1981). A new approach to decomposition of economic time series into permanent and transitory components with particular attention to measurement of the 'business cycle'. *Journal of Monetary Economics*, 7(2), 151–174. doi:10.1016/0304-3932(81)90040-4
- Bollerslev, T. (1986). Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3), 307–327. doi:10.1016/0304-4076(86)90063-1
- Carter, C. K., & Kohn, R. (1994). On Gibbs sampling for state space models. *Biometrika*, 81(3), 541–553. doi:10.1093/biomet/81.3.541
- Durbin, J., & Koopman, S. J. (2002). A simple and efficient simulation smoother for state space time series analysis. *Biometrika*, 89(3), 603–616. doi:10.1093/biomet/89.3.603
- Durbin, J., & Koopman, S. J. (2012). *Time Series Analysis by State Space Methods* (2nd ed.). Oxford University Press. doi:10.1093/acprof:oso/9780199641178.001.0001
- Gordon, N. J., Salmond, D. J., & Smith, A. F. M. (1993). Novel approach to nonlinear/non-Gaussian Bayesian state estimation. *IEE Proceedings F (Radar and Signal Processing)*, 140(2), 107–113. doi:10.1049/ip-f-2.1993.0015
- Kim, S., Shephard, N., & Chib, S. (1998). Stochastic volatility: Likelihood inference and comparison with ARCH models. *The Review of Economic Studies*, 65(3), 361–393. doi:10.1111/1467-937X.00050
- Koopman, S. J. (1997). Exact initial Kalman filtering and smoothing for nonstationary time series models. *Journal of the American Statistical Association*, 92(440), 1630–1638. doi:10.1080/01621459.1997.10473685
- Morley, J. C., Nelson, C. R., & Zivot, E. (2003). Why are the Beveridge–Nelson and unobserved-components decompositions of GDP so different? *The Review of Economics and Statistics*, 85(2), 235–243. doi:10.1162/003465303765299765
- Orphanides, A., & van Norden, S. (2002). The unreliability of output-gap estimates in real time. *The Review of Economics and Statistics*, 84(4), 569–583. doi:10.1162/003465302760556422
- Pitt, M. K., & Shephard, N. (1999). Filtering via simulation: Auxiliary particle filters. *Journal of the American Statistical Association*, 94(446), 590–599. doi:10.1080/01621459.1999.10474153
- Shephard, N., & Harvey, A. C. (1990). On the probability of estimating a deterministic component in the local level model. *Journal of Time Series Analysis*, 11(4), 339–347. doi:10.1111/j.1467-9892.1990.tb00062.x

References (2/2)

- Stock, J. H., & Watson, M. W. (1998). Median unbiased estimation of coefficient variance in a time-varying parameter model. *Journal of the American Statistical Association*, 93(441), 349–358. doi:10.1080/01621459.1998.10474116
- Stock, J. H., & Watson, M. W. (2007). Why has U.S. inflation become harder to forecast? *Journal of Money, Credit and Banking*, 39(s1), 3–33. doi:10.1111/j.1538-4616.2007.00014.x